

1. **Using the iris data set implement the KNN algorithm. Take different values for Test and training data set .Also use different values for k. Also find the accuracy level.**

CODE

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
dataset = pd.read_csv("iris.csv")
X = dataset.iloc[:, :-1].values
y = dataset.iloc[:, 4].values
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20)
from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors=5)
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)
from sklearn.metrics import classification_report, confusion_matrix
print(classification_report(y_test, y_pred))
from sklearn.metrics import accuracy_score
print ("Accuracy : ", accuracy_score(y_test, y_pred))
df = pd.DataFrame({'Real Values':y_test, 'Predicted Values':y_pred})
df
```

OUTPUT

```
              precision    recall  f1-score   support

   Setosa              1.00      1.00      1.00         9
  Versicolor          1.00      0.80      0.89        10
   Virginica          0.85      1.00      0.92        11

   accuracy              0.93              30
  macro avg              0.95      0.93      0.94        30
 weighted avg              0.94      0.93      0.93        30

Accuracy :  0.9333333333333333
```

Reference: <https://stackabuse.com/k-nearest-neighbors-algorithm-in-python-and-scikit-learn/>

2. **Download another data set suitable for the KNN and implement the KNN algorithm. Take different values for Test and training data set .Also use different values for k.**

CODE

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
dataset = pd.read_csv("cancer.csv")
dataset.head()
dataset.info()
```

```

X = dataset.iloc[:, 2:35].values
print(X)
y = dataset.iloc[:, 1].values
print(y)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20)
from sklearn.neighbors import KNeighborsClassifier
classifier = KNeighborsClassifier(n_neighbors=5)
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)
from sklearn.metrics import classification_report, confusion_matrix
print(classification_report(y_test, y_pred))
from sklearn.metrics import accuracy_score
print ("Accuracy : ", accuracy_score(y_test, y_pred))
df = pd.DataFrame({'Real Values':y_test, 'Predicted Values':y_pred})
df

```

OUTPUT

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 568 entries, 0 to 567
Data columns (total 32 columns):
#   Column      Non-Null Count  Dtype
---  -
0   842302      568 non-null    int64
1   M           568 non-null    object
2   17.99       568 non-null    float64
3   10.38       568 non-null    float64
4   122.8       568 non-null    float64
5   1001        568 non-null    float64
6   0.1184      568 non-null    float64
7   0.2776      568 non-null    float64
8   0.3001      568 non-null    float64
9   0.1471      568 non-null    float64
10  0.2419      568 non-null    float64
11  0.07871     568 non-null    float64
12  1.095       568 non-null    float64
13  0.9053      568 non-null    float64
14  0.500       568 non-null    float64

```

```

14  8.589      568 non-null    float64
15  153.4      568 non-null    float64
16  0.006399   568 non-null    float64
17  0.04904    568 non-null    float64
18  0.05373    568 non-null    float64
19  0.01587    568 non-null    float64
20  0.03003    568 non-null    float64
21  0.006193   568 non-null    float64
22  25.38      568 non-null    float64
23  17.33      568 non-null    float64
24  184.6      568 non-null    float64
25  2019       568 non-null    float64
26  0.1622     568 non-null    float64
27  0.6656     568 non-null    float64
28  0.7119     568 non-null    float64
29  0.2654     568 non-null    float64
30  0.4601     568 non-null    float64
31  0.1189     568 non-null    float64
dtypes: float64(30), int64(1), object(1)
memory usage: 142.1+ KB
memory usage: 142.1+ KB
[[2.057e+01 1.777e+01 1.329e+02 ... 1.860e-01 2.750e-01 8.902e-02]
 [1.969e+01 2.125e+01 1.300e+02 ... 2.430e-01 3.613e-01 8.758e-02]
 [1.142e+01 2.038e+01 7.758e+01 ... 2.575e-01 6.638e-01 1.730e-01]
 ...
 [1.660e+01 2.808e+01 1.083e+02 ... 1.418e-01 2.218e-01 7.820e-02]
 [2.060e+01 2.933e+01 1.401e+02 ... 2.650e-01 4.087e-01 1.240e-01]
 [7.760e+00 2.454e+01 4.792e+01 ... 0.000e+00 2.871e-01 7.039e-02]]
['M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M'
 'B' 'B' 'B' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M'
 'B' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'M' 'B' 'M' 'B' 'B' 'B' 'B' 'M' 'M'
 'B' 'M' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'M' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'M'
 'M' 'B' 'M' 'B' 'M' 'M' 'B' 'B' 'B' 'M' 'M' 'B' 'M' 'M' 'M' 'B' 'B' 'B'
 'M' 'B' 'B' 'M' 'M' 'B' 'B' 'B' 'M' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'B' 'M'
 'B' 'B' 'B' 'B' 'M' 'M' 'B' 'M' 'M' 'B' 'B' 'M' 'B' 'B' 'B' 'M' 'B' 'B'
 'M' 'B' 'M' 'B' 'B' 'B' 'B' 'B' 'B' 'B' 'B' 'M' 'B' 'B' 'B' 'B' 'B'
 'B' 'M' 'B' 'B' 'M' 'M' 'B' 'B' 'M' 'M' 'B' 'B' 'B' 'B' 'M' 'B' 'B' 'M'
 'M' 'M' 'B' 'M' 'B' 'M' 'B' 'B' 'B' 'M' 'B' 'B' 'M' 'M' 'B' 'M' 'M' 'M'
 'M' 'B' 'M' 'M' 'M' 'B' 'M' 'B' 'M' 'B' 'B' 'M' 'B' 'M' 'M' 'M' 'B'

```

	precision	recall	f1-score	support
B	0.96	0.93	0.94	71
M	0.89	0.93	0.91	43
accuracy			0.93	114
macro avg	0.92	0.93	0.93	114
weighted avg	0.93	0.93	0.93	114

Accuracy : 0.9298245614035088

Out[9]:

	Real Values	Predicted Values
0	B	B
1	B	B
2	B	B
3	B	B
4	B	B
...
109	B	B
110	B	B
111	M	M
112	M	M
113	B	B

114 rows × 2 columns

3. Using iris data set, implement naive bayes classification for different naive Bayes classification algorithms.((i) gaussian (ii) bernoulli etc)

- Find out the accuracy level w.r.t to each algorithm
- Display the no:of mislabeled classification from test data set
- List out the class labels of the mismatching records

i)CODE

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
dataset = pd.read_csv('iris.csv')
X = dataset.iloc[:, :4].values
```

```
y = dataset['variety'].values
dataset.head(5)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
from sklearn.naive_bayes import GaussianNB
classifier = GaussianNB()
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)
y_pred
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
from sklearn.metrics import accuracy_score
print ("Accuracy : ", accuracy_score(y_test, y_pred))
cm
df = pd.DataFrame({'Real Values':y_test, 'Predicted Values':y_pred})
df
```

OUTPUT

Accuracy : 1.0

:

	Real Values	Predicted Values
0	Versicolor	Versicolor
1	Versicolor	Versicolor
2	Setosa	Setosa
3	Setosa	Setosa
4	Setosa	Setosa
5	Setosa	Setosa
6	Setosa	Setosa
7	Versicolor	Versicolor
8	Setosa	Setosa
9	Versicolor	Versicolor
10	Setosa	Setosa
11	Setosa	Setosa
12	Setosa	Setosa
13	Setosa	Setosa
14	Virginica	Virginica
15	Setosa	Setosa
16	Versicolor	Versicolor
17	Versicolor	Versicolor
18	Virginica	Virginica
19	Virginica	Virginica
20	Setosa	Setosa
21	Setosa	Setosa
22	Versicolor	Versicolor
23	Versicolor	Versicolor
24	Versicolor	Versicolor
25	Virginica	Virginica
26	Versicolor	Versicolor
27	Versicolor	Versicolor
28	Virginica	Virginica
29	Setosa	Setosa

CODE

```

import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
dataset = pd.read_csv('iris.csv')
X = dataset.iloc[:,4].values
y = dataset['variety'].values
dataset.head(5)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
from sklearn.naive_bayes import BernoulliNB
classifier = BernoulliNB()
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)
y_pred
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
from sklearn.metrics import accuracy_score
print ("Accuracy : ", accuracy_score(y_test, y_pred))
cm
# df = pd.DataFrame({'Real Values':y_test, 'Predicted Values':y_pred})
# df

```

OUTPUT

```

Accuracy :  0.23333333333333334
array([[ 7,  0,  0],
       [13,  0,  0],
       [10,  0,  0]], dtype=int64)

```

ii)CODE

```

import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

```

```
dataset = pd.read_csv('iris.csv')
X = dataset.iloc[:, :4].values
y = dataset['variety'].values
dataset.head(5)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
from sklearn.naive_bayes import BernoulliNB
classifier = BernoulliNB()
classifier.fit(X_train, y_train)
y_pred = classifier.predict(X_test)
y_pred
#from sklearn.metrics import confusion_matrix
#cm = confusion_matrix(y_test, y_pred)
from sklearn.metrics import accuracy_score
print ("Accuracy : ", accuracy_score(y_test, y_pred))
#cm
df = pd.DataFrame({'Real Values':y_test, 'Predicted Values':y_pred})
df
```

OUTPUT

Accuracy : 0.2666666666666666

	Real Values	Predicted Values
0	Virginica	Versicolor
1	Virginica	Versicolor
2	Virginica	Versicolor
3	Versicolor	Versicolor
4	Setosa	Versicolor
5	Setosa	Versicolor
6	Setosa	Versicolor
7	Setosa	Versicolor
8	Setosa	Versicolor
9	Setosa	Versicolor
10	Versicolor	Versicolor
11	Setosa	Versicolor
12	Setosa	Versicolor
13	Setosa	Versicolor
14	Versicolor	Versicolor
15	Setosa	Versicolor

16	Versicolor	Versicolor
17	Virginica	Versicolor
18	Virginica	Versicolor
19	Versicolor	Versicolor
20	Virginica	Versicolor
21	Versicolor	Versicolor
22	Virginica	Versicolor
23	Virginica	Versicolor
24	Versicolor	Versicolor
25	Setosa	Versicolor
26	Versicolor	Versicolor
27	Virginica	Versicolor
28	Setosa	Versicolor
29	Setosa	Versicolor

References:

<https://towardsdatascience.com/machine-learning-basics-naive-bayes-classification-964af6f2a965>

https://scikit-learn.org/stable/modules/classes.html#module-sklearn.naive_bayes

4. Use car details CSV file and implement decision tree algorithm

- Find out the accuracy level.
- Display the no:of mislabeled classification from test data set
- List out the class labels of the mismatching records

CODE

```
import os
```

```
import numpy as np
```

```
import pandas as pd
```

```
import numpy as np, pandas as pd
```

```
import matplotlib.pyplot as plt
```

```
from sklearn import tree, metrics, model_selection
```

```
data =
```

```
pd.read_csv('car.csv',names=['buying','maint','doors','persons','lug_boot','safety',  
, 'class'])
```

```
data.head()
```

	buying	maint	doors	persons	lug_boot	safety	class
0	vhigh	vhigh	2	2	small	low	unacc
1	vhigh	vhigh	2	2	small	med	unacc
2	vhigh	vhigh	2	2	small	high	unacc
3	vhigh	vhigh	2	2	med	low	unacc
4	vhigh	vhigh	2	2	med	med	unacc

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1728 entries, 0 to 1727
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   buying      1728 non-null   object
1   maint       1728 non-null   object
2   doors       1728 non-null   object
3   persons     1728 non-null   object
4   lug_boot    1728 non-null   object
5   safety      1728 non-null   object
6   class       1728 non-null   object
dtypes: object(7)
memory usage: 94.6+ KB
```

```
data['class'],class_names = pd.factorize(data['class'])
```

```
print(class_names)
```

```
print(data['class'].unique())
```

```
Index(['unacc', 'acc', 'vgood', 'good'], dtype='object')
[0 1 2 3]
```

```
data['buying'],_ = pd.factorize(data['buying'])
```

```
data['maint'],_ = pd.factorize(data['maint'])
```

```
data['doors'],_ = pd.factorize(data['doors'])
```

```
data['persons'],_ = pd.factorize(data['persons'])
```

```
data['lug_boot'],_ = pd.factorize(data['lug_boot'])
```

```
data['safety'],_ = pd.factorize(data['safety'])
```

```
data.head()
```

	buying	maint	doors	persons	lug_boot	safety	class
0	0	0	0	0	0	0	0
1	0	0	0	0	0	1	0
2	0	0	0	0	0	2	0
3	0	0	0	0	1	0	0
4	0	0	0	0	1	1	0

data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1728 entries, 0 to 1727
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   buying      1728 non-null   int64
1   maint       1728 non-null   int64
2   doors       1728 non-null   int64
3   persons     1728 non-null   int64
4   lug_boot    1728 non-null   int64
5   safety      1728 non-null   int64
6   class       1728 non-null   int64
dtypes: int64(7)
memory usage: 94.6 KB
```

```
X = data.iloc[:, :-1]
```

```
y = data.iloc[:, -1]
```

```
# split data randomly into 70% training and 30% test
```

```
X_train, X_test, y_train, y_test = model_selection.train_test_split(X, y,
test_size=0.3, random_state=0)
```

```
# train the decision tree
```

```
dtree = tree.DecisionTreeClassifier(criterion='entropy', max_depth=3,  
random_state=0)  
dtree.fit(X_train, y_train)
```

```
DecisionTreeClassifier(criterion='entropy', max_depth=3, random_state=0)
```

```
# use the model to make predictions with the test data  
y_pred = dtree.predict(X_test)  
# how did our model perform?  
accuracy = metrics.accuracy_score(y_test, y_pred)  
print('Accuracy: {:.2f}'.format(accuracy))
```

Accuracy: 0.82

```
count_misclassified = (y_test != y_pred).sum()  
print('Misclassified samples: {}'.format(count_misclassified))
```

Misclassified samples: 96

References:

<https://www.24tutorials.com/machine-learning/case-study-decision-tree-model-for-car-quality/>
https://notebook.community/bMzi/ML_in_Finance/0210_DecisionTrees
<https://stackabuse.com/decision-trees-in-python-with-scikit-learn/>

For Data Sets Refer:

<https://www.kaggle.com/> (for data set)
<http://archive.ics.uci.edu/ml/datasets.php>

5.Implement Simple and multiple linear regression for the data sets 'student_score.csv' and 'company_data .csv' respectively

Ref: <https://stackabuse.com/linear-regression-in-python-with-scikit-learn/>

```
1) import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
#data set contains details of no.of hours spend by students for studt and their marks
student = pd.read_csv('student_scores.csv')
student.head()
```

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30

```
2) student.describe()
```


	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

```
3) student.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25 entries, 0 to 24
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Hours   25 non-null        float64
1   Scores  25 non-null        int64
dtypes: float64(1), int64(1)
memory usage: 528.0 bytes
```

```
4) import matplotlib.pyplot as plt
Xax=student.iloc[:,0]
Yax=student.iloc[:,1]
plt.scatter(Xax,Yax)
plt.xlabel("No.of hours")
plt.ylabel("Score")
plt.title("Student scores")
```

plt.show()



5) #perform the simple linear regression model

#Equation : $Y = w_0 + w_1 \cdot x$

#Here $Y(\text{marks}) = w_0 + w_1 \cdot x$

#create x as hours and Y as marks

$X = \text{student.iloc[:, :-1]}$

$y = \text{student.iloc[:, 1]}$

$\text{print}(X)$

	Hours
0	2.5
1	5.1
2	3.2
3	8.5
4	3.5
5	1.5
6	9.2
7	5.5
8	8.3
9	2.7
10	7.7
11	5.9
12	4.5
13	3.3
14	1.1
15	8.9
16	2.5
17	1.9
18	6.1
19	7.4
20	2.7
21	4.8
22	3.8
23	6.9
24	7.8

6) `print(y)`

0	21
1	47
2	27
3	75
4	30
5	20
6	88
7	60
8	81
9	25
10	85
11	62
12	41
13	42
14	17
15	95
16	30
17	24
18	67
19	69
20	30
21	54
22	35
23	76
24	86

Name: Scores, dtype: int64


```
7) from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
print(X_train)
```

```
Hours
2      3.2
22     3.8
1      5.1
6      9.2
19     7.4
7      5.5
21     4.8
16     2.5
9      2.7
17     1.9
8      8.3
13     3.3
20     2.7
3      8.5
18     6.1
15     8.9
11     5.9
0      2.5
5      1.5
24     7.8
```

```
8) from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train, y_train)
```

```
LinearRegression()
```

```
9) print(regressor.intercept_)
```

```
3.4157789720979608
```

```
10) print(regressor.coef_)
```

```
print(regressor.coef_)
```

```
[9.544138]
```

```
11) y_pred = regressor.predict(X_test)
for(i,j) in zip(y_test,y_pred):
    if i!=j:
        print("Actual value :",i,"Predicted value :",j)
print("Number of mislabeled points from test data set :", (y_test != y_pred).sum())
```

```

Actual value : 76 Predicted value : 69.2703311556657
Actual value : 41 Predicted value : 46.36439996138126
Actual value : 85 Predicted value : 76.90564155376049
Actual value : 17 Predicted value : 13.914330769478324
Actual value : 30 Predicted value : 36.82026196376275
Number of mislabeled points from test data set : 5

```

```

12) from sklearn import metrics
    print("Mean Absolute error :", metrics.mean_absolute_error(y_test,y_pred))
    print("Mean Squared error :", metrics.mean_squared_error(y_test,y_pred))
    print("Root Mean Squared error :",
          np.sqrt(metrics.mean_squared_error(y_test,y_pred)))

```

```

Mean Absolute error : 6.0188716892478995
Mean Squared error : 39.124239242163426
Root Mean Squared error : 6.254937189306014

```


```

13) import matplotlib.pyplot as plt
    c=X_test['Hours'].count()
    xax=np.arange(c)
    print(xax)
    X_axis = np.arange(len(xax))
    plt.bar(X_axis-0.2, y_test, 0.6, label='Actual')
    plt.bar(X_axis+0.2, y_pred, 0.6, label='Predicted')

    plt.xlabel("Test Records")
    plt.ylabel("Marks")
    plt.title("Student Score prediction")
    plt.legend()
    plt.show()

```

```
[0 1 2 3 4]
```



Multiple Regression

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
advertising = pd.read_csv('Company_data.csv')
advertising.head()
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

```
advertising.describe()
```


	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

```
advertising.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    TV          200 non-null    float64
1    Radio       200 non-null    float64
2    Newspaper   200 non-null    float64
3    Sales       200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

```
import matplotlib.pyplot as plt
import seaborn as sns
sns.pairplot(advertising, x_vars=['TV', 'Radio', 'Newspaper'],
             y_vars='Sales', height=5, aspect=1, kind='scatter')
```

```
plt.show()
```



```
#perform the multiple linear regression model
#Equation :  $Y = w_0 + w_1.x_1 + w_2.x_2 + w_3.x_3$ 
#Here  $Y(\text{sales}) = w_0 + w_1.x_1(\text{TV}) + w_2.x_2(\text{Radio}) + w_3.x_3(\text{Newspaper})$ 
#create x and Y as sales
```

```
X = advertising.iloc[:, :-1]
print(X)
```

	TV	Radio	Newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4
...
195	38.2	3.7	13.8
196	94.2	4.9	8.1
197	177.0	9.3	6.4
198	283.6	42.0	66.2
199	232.1	8.6	8.7

```
[200 rows x 3 columns]
```

```
y = advertising.iloc[:, -1]
print(y)
```

0	22.1
1	10.4
2	12.0
3	16.5
4	17.9
...	...
195	7.6
196	14.0
197	14.8
198	25.5
199	18.4

```
Name: Sales, Length: 200, dtype: float64
```

```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
print(X_train)

```

	TV	Radio	Newspaper
74	213.4	24.6	13.1
108	13.1	0.4	25.6
24	62.3	12.6	18.3
185	205.0	45.1	19.6
70	199.1	30.6	38.7
..
132	8.4	27.2	2.1
72	26.8	33.0	19.3
91	28.6	1.5	33.0
147	243.2	49.0	44.3
78	5.4	29.9	9.4

```
[140 rows x 3 columns]
```

```

from sklearn.linear_model import LinearRegression
regressor = LinearRegression()
regressor.fit(X_train, y_train)

```

```
LinearRegression()
```

```
print(regressor.intercept_)
```

```
4.834885358384083
```

```
print(regressor.coef_)
```

```
[ 0.05441868  0.11045214 -0.00292727]
```

```

y_pred = regressor.predict(X_test)
for(i,j) in zip(y_test,y_pred):
    if i!=j:
        print("Actual value :",i,"Predicted value :",j)
print("Number of mislabeled points from test data set :", (y_test != y_pred).sum())

```

Actual value :	18.0	Predicted value :	19.6231329850721
Actual value :	13.3	Predicted value :	13.308438358662285
Actual value :	11.6	Predicted value :	12.629676247192208
Actual value :	17.0	Predicted value :	19.47105688754111
Actual value :	12.3	Predicted value :	11.590783196229005
Actual value :	10.7	Predicted value :	11.348323111604875
Actual value :	11.5	Predicted value :	12.193874211599226
Actual value :	1.6	Predicted value :	9.221415842312133
Actual value :	14.2	Predicted value :	14.498654917672432
Actual value :	23.2	Predicted value :	22.4195357567519
Actual value :	12.6	Predicted value :	12.633285935460439
Actual value :	13.2	Predicted value :	13.506916188363006
Actual value :	18.9	Predicted value :	21.30047124135335
Actual value :	11.3	Predicted value :	9.614344565789548
Actual value :	17.3	Predicted value :	16.895894130341553
Actual value :	17.7	Predicted value :	19.461157758162322
Actual value :	20.2	Predicted value :	22.265150396440184
Actual value :	10.3	Predicted value :	12.653339706591243
Actual value :	16.6	Predicted value :	18.225248894921528
Actual value :	18.0	Predicted value :	17.77531629354258
Actual value :	18.4	Predicted value :	18.389881966359212
Actual value :	6.7	Predicted value :	7.120487442072934
Actual value :	22.1	Predicted value :	21.329147279261107
Actual value :	7.0	Predicted value :	8.107233580890917
Actual value :	17.2	Predicted value :	17.241601167641473
Actual value :	15.6	Predicted value :	15.932854944685918
Actual value :	9.2	Predicted value :	9.814012958114937
Actual value :	7.6	Predicted value :	8.07181587152485
Actual value :	16.6	Predicted value :	17.24408925929915
Actual value :	16.1	Predicted value :	21.53718441018264

```
Actual value : 21.8 Predicted value : 22.11714873643335
Actual value : 12.9 Predicted value : 12.24172273842917
Actual value : 17.5 Predicted value : 15.956455247767064
Actual value : 16.7 Predicted value : 14.855431394128658
Actual value : 21.2 Predicted value : 19.70279428720466
Actual value : 20.8 Predicted value : 22.746586722943107
Actual value : 7.3 Predicted value : 9.414955243763849
Actual value : 19.8 Predicted value : 18.797882366422925
Actual value : 12.9 Predicted value : 13.932861260132384
Actual value : 19.8 Predicted value : 21.68847856864288
Actual value : 10.8 Predicted value : 11.507027883270196
Actual value : 19.8 Predicted value : 21.135922336380666
Actual value : 16.4 Predicted value : 16.13142408067001
Actual value : 20.1 Predicted value : 20.39027162911635
Actual value : 10.4 Predicted value : 11.055958359288084
Actual value : 15.2 Predicted value : 15.319170605692749
Actual value : 16.8 Predicted value : 20.07719756905186
Actual value : 15.9 Predicted value : 15.519157553678733
Actual value : 11.6 Predicted value : 12.2047855314722
Actual value : 18.9 Predicted value : 20.759113647419312
Actual value : 17.6 Predicted value : 14.899620097103586
Actual value : 16.6 Predicted value : 18.119983975440608
Actual value : 19.9 Predicted value : 17.01161015049727
Actual value : 17.9 Predicted value : 16.944628168884364
Actual value : 12.6 Predicted value : 9.001742559147871
Actual value : 7.2 Predicted value : 10.489892255018589
Actual value : 10.3 Predicted value : 12.580375265443143
Number of mislabeled points from test data set : 60
```


```
from sklearn import metrics
print("Mean Absolute error :", metrics.mean_absolute_error(y_test,y_pred))
print("Mean Squared error :", metrics.mean_squared_error(y_test,y_pred))
print("Root Mean Squared error :", np.sqrt(metrics.mean_squared_error(y_test,y_pred)))
```

```
Mean Absolute error : 1.4439815071936666
Mean Squared error : 3.9021227595995556
Root Mean Squared error : 1.9753791432531518
```

```
import matplotlib.pyplot as plt
c=X_test['TV'].count()
xax=np.arange(c)
print(xax)
X_axis = np.arange(len(xax))
plt.bar(X_axis-0.2, y_test, 0.6, label='Actual')
plt.bar(X_axis+0.2, y_pred, 0.6, label='Predicted')

plt.xlabel("Sales")
plt.ylabel("Actual/Predicted")
plt.title("Sales prediction")
plt.legend()
plt.show()
```

```
[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47
48 49 50 51 52 53 54 55 56 57 58 59]
```



Neural Networks

1.Create a neural network for the given ‘houseprice.csv’ to predict the whether price of the house is above or below median value or not.

```
import tensorflow as tf
```

```
import keras
```

```
import pandas
```

```
import sklearn
```

```
import matplotlib
```

```
import pandas as pd
```

```
df=pd.read_csv('housepricedata.csv')
```

```
df
```

	LotArea	OverallQual	OverallCond	TotalBsmntSF	FullBath	HalfBath	BedroomAbvGr	TotRmsAbvGrd	Fireplaces	GarageArea	AboveMedianPrice
0	8450	7	5	856	2	1	3	8	0	548	1
1	9600	6	8	1262	2	0	3	6	1	460	1
2	11250	7	5	920	2	1	3	6	1	608	1
3	9550	7	5	756	1	0	3	7	1	642	0
4	14260	8	5	1145	2	1	4	9	1	836	1
...
1455	7917	6	5	953	2	1	3	7	1	460	1
1456	13175	6	6	1542	2	0	3	7	2	500	1
1457	9042	7	9	1152	2	0	4	9	2	252	1
1458	9717	5	6	1078	1	0	2	5	0	240	0
1459	9937	5	6	1256	1	1	3	6	0	276	0

1460 rows × 11 columns


```
dataset = df.values
```

```
dataset
```

```
array([[ 8450,    7,    5, ...,    0,   548,    1],
       [ 9600,    6,    8, ...,    1,   460,    1],
       [11250,    7,    5, ...,    1,   608,    1],
       ...,
       [ 9042,    7,    9, ...,    2,   252,    1],
       [ 9717,    5,    6, ...,    0,   240,    0],
       [ 9937,    5,    6, ...,    0,   276,    0]], dtype=int64)
```

```
X = dataset[:,0:10]
```

```
Y = dataset[:,10]
```

```
from sklearn import preprocessing
```

```
min_max_scaler = preprocessing.MinMaxScaler()
```

```
X_scale = min_max_scaler.fit_transform(X)
```

```
X_scale
```

```
array([[0.0334198 , 0.66666667, 0.5      , ..., 0.5      , 0.      ,
        0.3864598 ],
       [0.03879502, 0.55555556, 0.875    , ..., 0.33333333, 0.33333333,
        0.32440056],
       [0.04650728, 0.66666667, 0.5      , ..., 0.33333333, 0.33333333,
        0.42877292],
       ...,
       [0.03618687, 0.66666667, 1.      , ..., 0.58333333, 0.66666667,
        0.17771509],
       [0.03934189, 0.44444444, 0.625    , ..., 0.25      , 0.      ,
        0.16925247],
       [0.04037019, 0.44444444, 0.625    , ..., 0.33333333, 0.      ,
        0.19464034]])
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_val_and_test, Y_train, Y_val_and_test = train_test_split(X_scale, Y,
test_size=0.3)
```

```
X_val, X_test, Y_val, Y_test = train_test_split(X_val_and_test, Y_val_and_test,
test_size=0.5)
```

```
print(X_train.shape, X_val.shape, X_test.shape, Y_train.shape, Y_val.shape, Y_test.shape)
```

```
(1022, 10) (219, 10) (219, 10) (1022,) (219,) (219,)
```

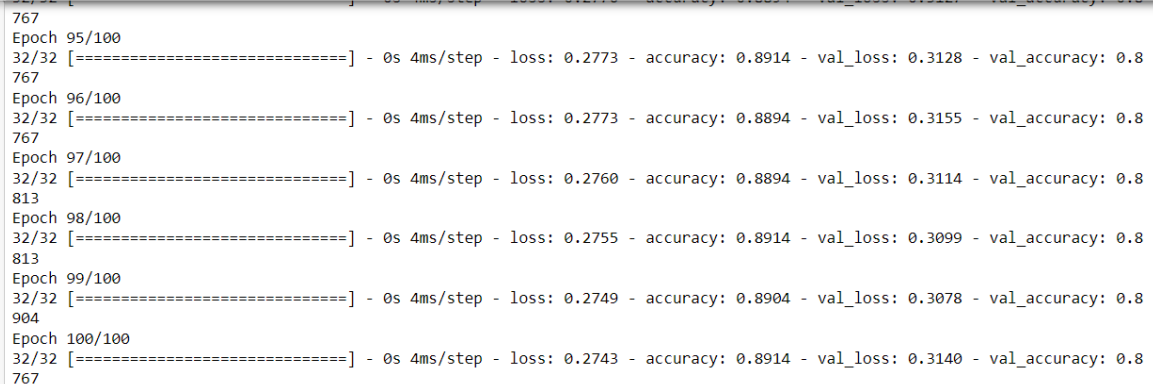
```
from keras.models import Sequential
```

```
from keras.layers import Dense
```

```
model = Sequential([  
    Dense(32, activation='relu', input_shape=(10,)),  
    Dense(32, activation='relu'),  
    Dense(1, activation='sigmoid'),  
])
```

```
model.compile(optimizer='sgd',  
              loss='binary_crossentropy',  
              metrics=['accuracy'])
```

```
hist = model.fit(X_train, Y_train,  
                 batch_size=32, epochs=100,  
                 validation_data=(X_val, Y_val))
```



```
767  
Epoch 95/100  
32/32 [=====] - 0s 4ms/step - loss: 0.2773 - accuracy: 0.8914 - val_loss: 0.3128 - val_accuracy: 0.8  
767  
Epoch 96/100  
32/32 [=====] - 0s 4ms/step - loss: 0.2773 - accuracy: 0.8894 - val_loss: 0.3155 - val_accuracy: 0.8  
767  
Epoch 97/100  
32/32 [=====] - 0s 4ms/step - loss: 0.2760 - accuracy: 0.8894 - val_loss: 0.3114 - val_accuracy: 0.8  
813  
Epoch 98/100  
32/32 [=====] - 0s 4ms/step - loss: 0.2755 - accuracy: 0.8914 - val_loss: 0.3099 - val_accuracy: 0.8  
813  
Epoch 99/100  
32/32 [=====] - 0s 4ms/step - loss: 0.2749 - accuracy: 0.8904 - val_loss: 0.3078 - val_accuracy: 0.8  
904  
Epoch 100/100  
32/32 [=====] - 0s 4ms/step - loss: 0.2743 - accuracy: 0.8914 - val_loss: 0.3140 - val_accuracy: 0.8  
767
```

```
model.evaluate(X_test, Y_test)[1]
```

```
7/7 [=====] - 0s 3ms/step - loss: 0.2653 - accuracy: 0.8584  
0.8584474921226501
```

```
import matplotlib.pyplot as plt
```

```
plt.plot(hist.history['loss'])
```

```
plt.plot(hist.history['val_loss'])
```


```
plt.title('Model loss')
```

```
plt.ylabel('Loss')
```

```
plt.xlabel('Epoch')
```

```
plt.legend(['Train', 'Val'], loc='upper right')
```

```
plt.show()
```




```
plt.plot(hist.history['accuracy'])
```

```
plt.plot(hist.history['val_accuracy'])
```

```
plt.title('Model accuracy')
```

```
plt.ylabel('Accuracy')
```

```
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='lower right')
plt.show()
```



```
model_2 = Sequential([
    Dense(1000, activation='relu', input_shape=(10,)),
    Dense(1000, activation='relu'),
    Dense(1000, activation='relu'),
    Dense(1000, activation='relu'),
    Dense(1, activation='sigmoid'),
])
model_2.compile(optimizer='adam',
                loss='binary_crossentropy',
                metrics=['accuracy'])
hist_2 = model_2.fit(X_train, Y_train,
                    batch_size=32, epochs=100,
                    validation_data=(X_val, Y_val))
```

```


8721
Epoch 95/100
32/32 [=====] - 1s 38ms/step - loss: 0.2040 - accuracy: 0.9129 - val_loss: 0.4571 - val_accuracy: 0.
8904
Epoch 96/100
32/32 [=====] - 1s 42ms/step - loss: 0.1727 - accuracy: 0.9315 - val_loss: 0.4910 - val_accuracy: 0.
8950
Epoch 97/100
32/32 [=====] - 1s 45ms/step - loss: 0.1240 - accuracy: 0.9472 - val_loss: 0.6341 - val_accuracy: 0.
8813
Epoch 98/100
32/32 [=====] - 1s 37ms/step - loss: 0.1060 - accuracy: 0.9569 - val_loss: 0.6325 - val_accuracy: 0.
8767
Epoch 99/100
32/32 [=====] - 1s 39ms/step - loss: 0.1429 - accuracy: 0.9442 - val_loss: 0.7589 - val_accuracy: 0.
8676
Epoch 100/100
32/32 [=====] - 1s 39ms/step - loss: 0.1572 - accuracy: 0.9315 - val_loss: 0.7380 - val_accuracy: 0.
8676

```

```

plt.plot(hist_2.history['loss'])
plt.plot(hist_2.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper right')
plt.show()

```




```

plt.plot(hist_2.history['accuracy'])

```

```
plt.plot(hist_2.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='lower right')
plt.show()
```



```
from keras.layers import Dropout
from keras import regularizers
```

```
model_3 = Sequential([
    Dense(1000, activation='relu', kernel_regularizer=regularizers.l2(0.01),
input_shape=(10,)),
    Dropout(0.3),
    Dense(1000, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    Dropout(0.3),
    Dense(1000, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    Dropout(0.3),
```

```

Dense(1000, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
Dropout(0.3),
Dense(1, activation='sigmoid', kernel_regularizer=regularizers.l2(0.01)),
])

```

```

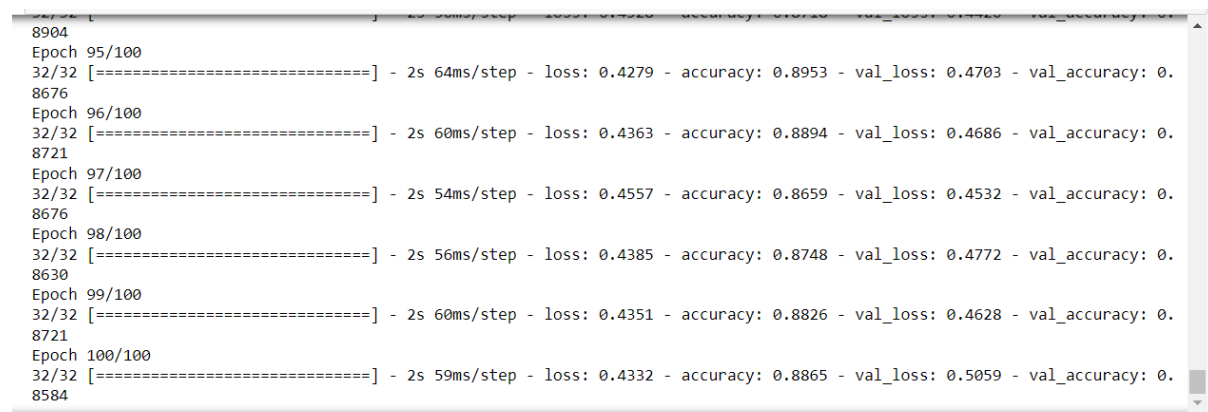
model_3.compile(optimizer='adam',
                loss='binary_crossentropy',
                metrics=['accuracy'])

```

```

hist_3 = model_3.fit(X_train, Y_train,
                    batch_size=32, epochs=100,
                    validation_data=(X_val, Y_val))

```



```


8904
Epoch 95/100
32/32 [=====] - 2s 64ms/step - loss: 0.4279 - accuracy: 0.8953 - val_loss: 0.4703 - val_accuracy: 0.
8676
Epoch 96/100
32/32 [=====] - 2s 60ms/step - loss: 0.4363 - accuracy: 0.8894 - val_loss: 0.4686 - val_accuracy: 0.
8721
Epoch 97/100
32/32 [=====] - 2s 54ms/step - loss: 0.4557 - accuracy: 0.8659 - val_loss: 0.4532 - val_accuracy: 0.
8676
Epoch 98/100
32/32 [=====] - 2s 56ms/step - loss: 0.4385 - accuracy: 0.8748 - val_loss: 0.4772 - val_accuracy: 0.
8630
Epoch 99/100
32/32 [=====] - 2s 60ms/step - loss: 0.4351 - accuracy: 0.8826 - val_loss: 0.4628 - val_accuracy: 0.
8721
Epoch 100/100
32/32 [=====] - 2s 59ms/step - loss: 0.4332 - accuracy: 0.8865 - val_loss: 0.5059 - val_accuracy: 0.
8584

```

```

plt.plot(hist_3.history['loss'])
plt.plot(hist_3.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='upper right')
plt.ylim(top=1.2, bottom=0)
plt.show()

```



```
plt.plot(hist_3.history['accuracy'])
plt.plot(hist_3.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Val'], loc='lower right')
plt.show()
```